

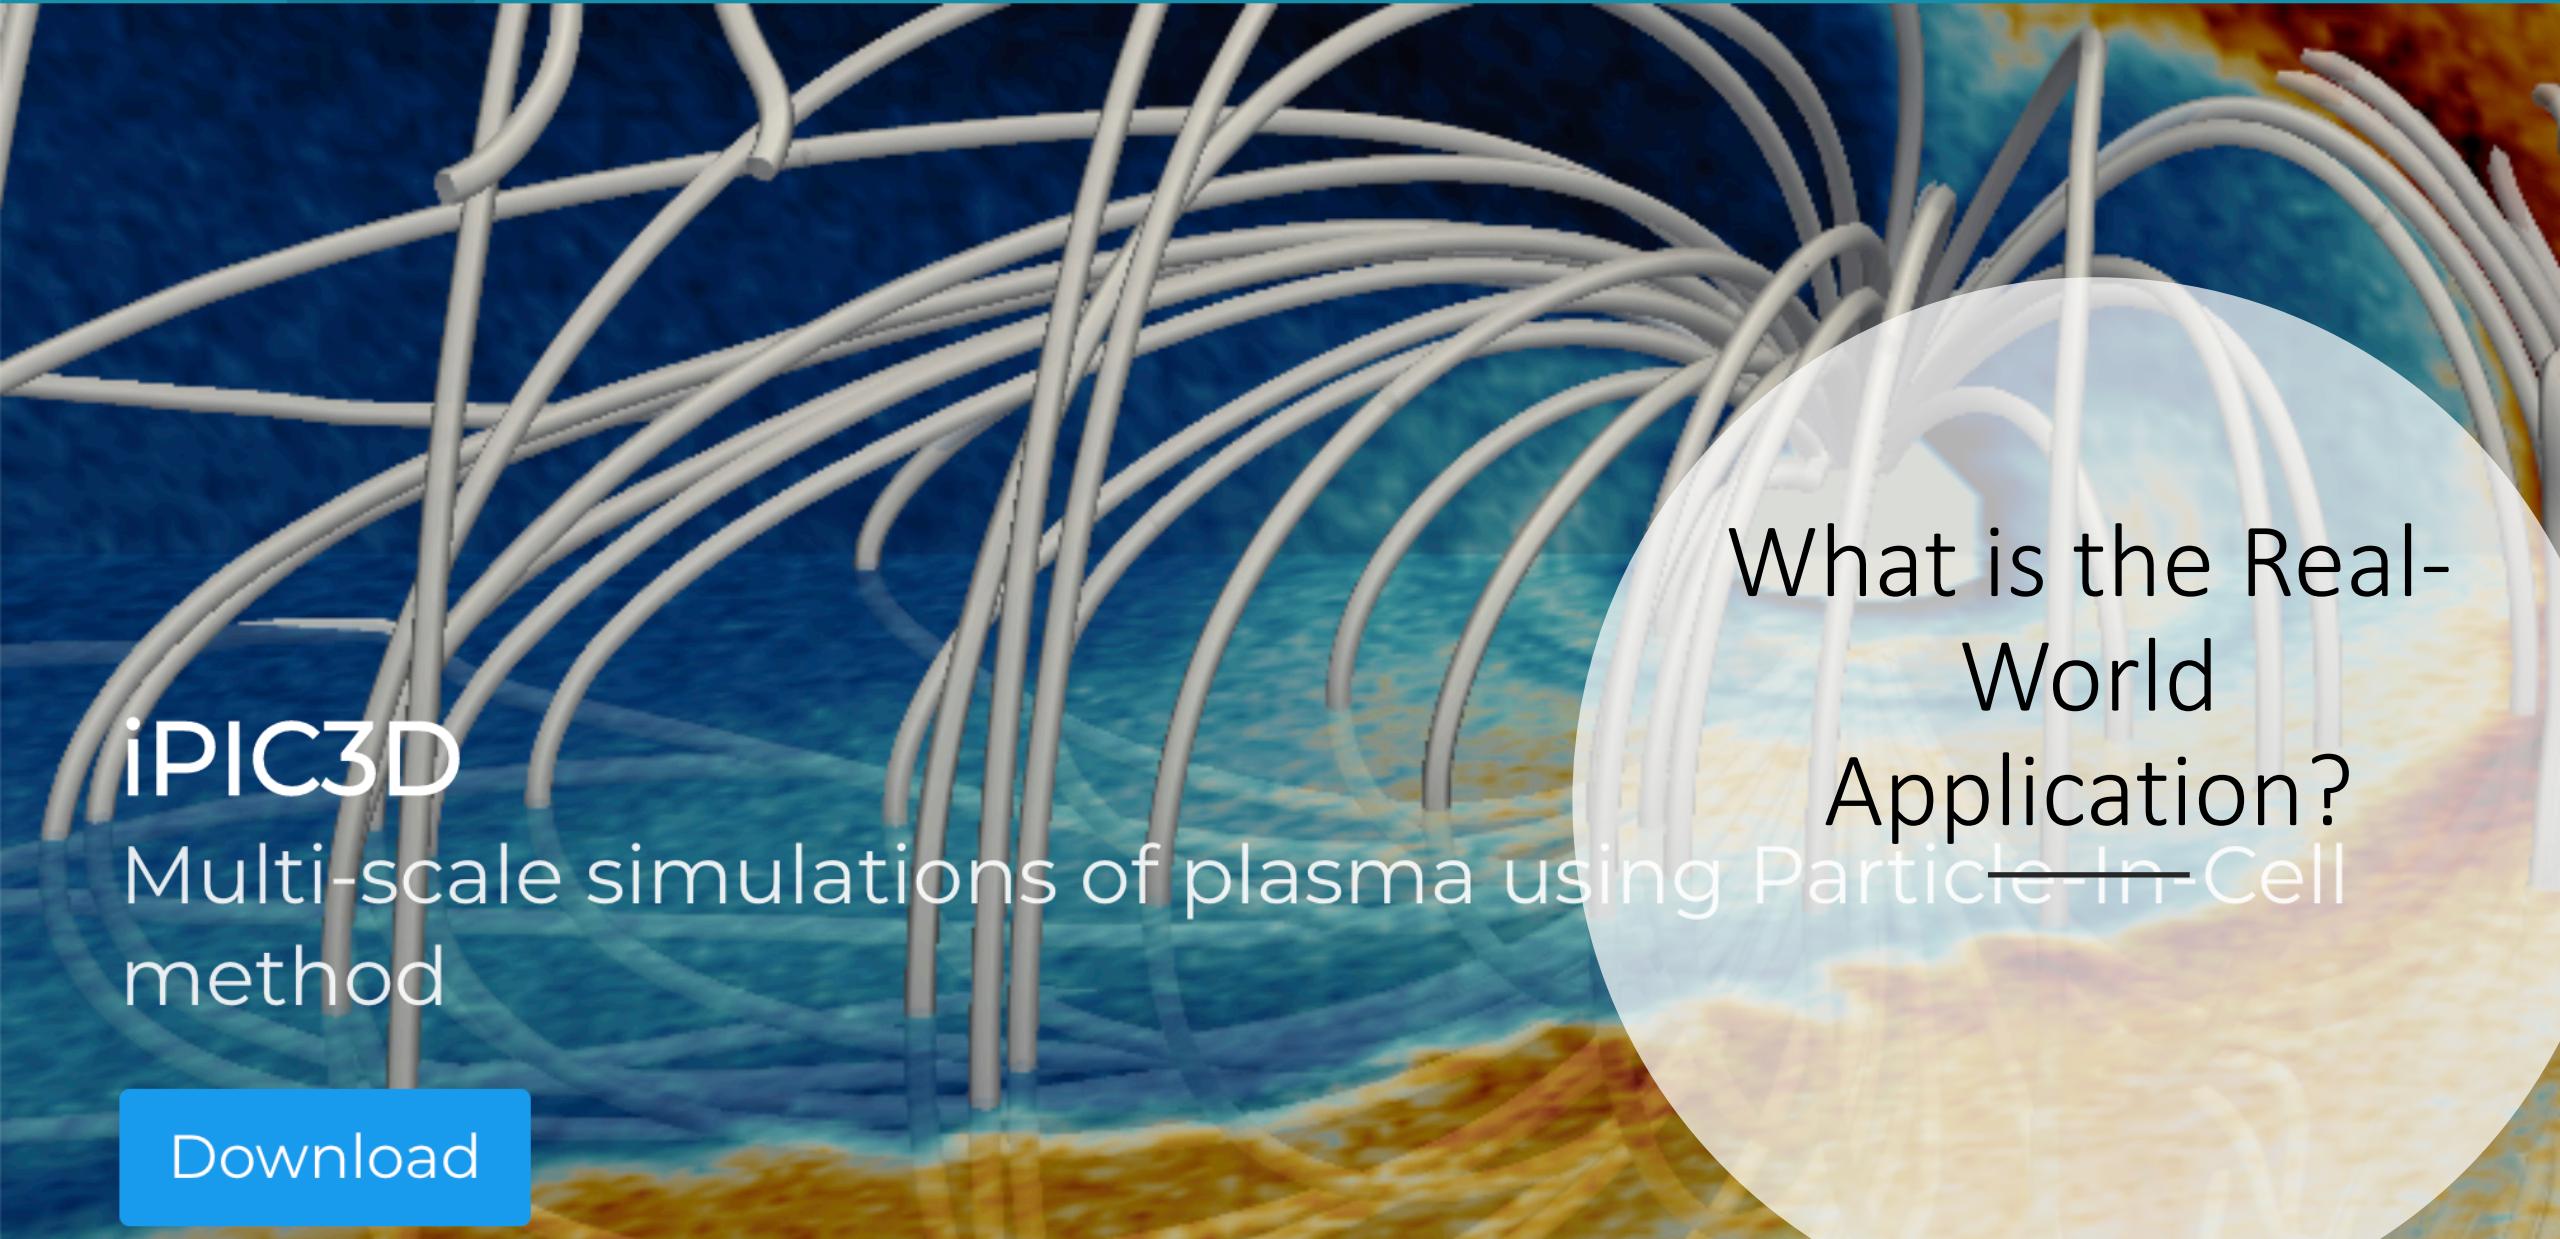
DD2360 Final Project

Stefano Markidis



Objective

- Take a real-world scientific code and
 - Port it to GPU with CUDA
 - Measure and analyze performance
 - Compare Performance with CPU and OpenACC (only A grade)
 - Optimize GPU code
 - Evaluate performance improvement



iPIC3D
Multi-scale simulations of plasma using Particle-In-Cell method

Download

What is the Real-
World
Application?

sputniPIC – A Simplified Serial C iPIC3D



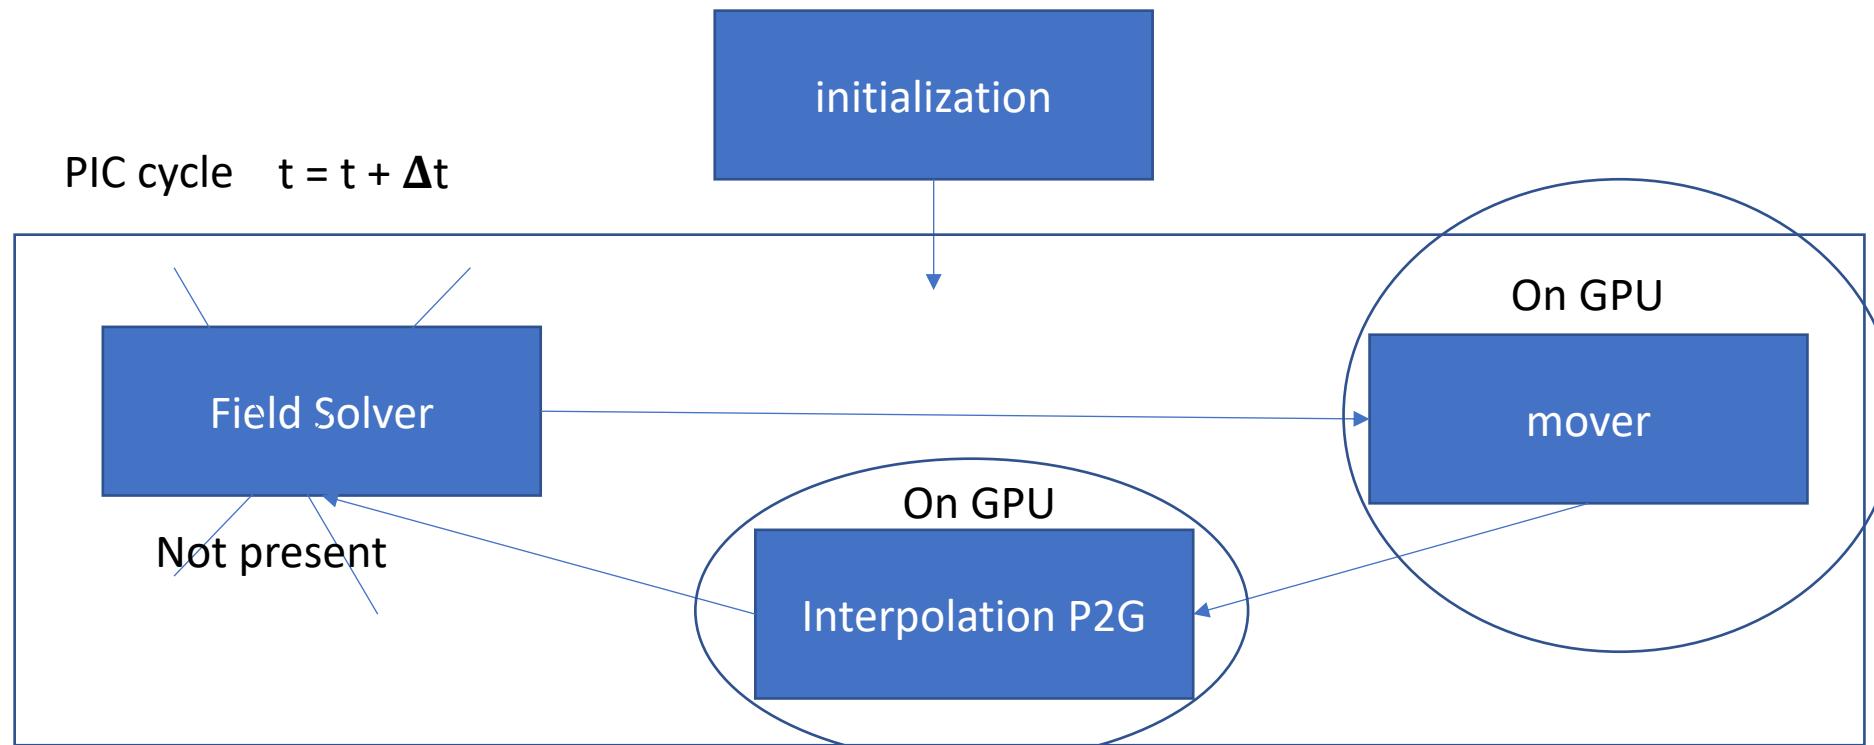
<https://github.com/iPIC3D/sputniPIC>



We will port and optimize two
functions of the sputniPIC code

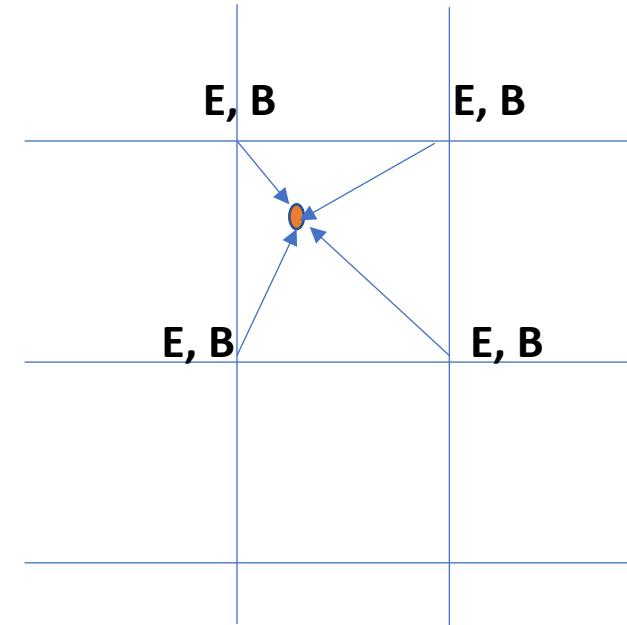
What is a Particle-in-Cell (PIC) Method?

- A numerical technique for simulating plasmas
 - Electron and protons are computational particles
 - At each simulation step, particles are moved solving an ODE for each particle
 - At each simulation step, particles deposit charge and weight on a grid
- In this project, we focus on two steps of the PIC method



Particle Mover

- Solve two ODEs for each particle
 - $\mathbf{v} = \mathbf{v} + q/m (\mathbf{E} + \mathbf{v} \times \mathbf{B})dt$
 - $\mathbf{x} = \mathbf{x} + \mathbf{v} dt$
- E and B are defined only grid points and not at the particle position
 - Need to interpolate (linearly) E and B at the particle position
- In Particle.cu, mover_PC
 - PC = use a predictor corrector scheme



How Many Particles?

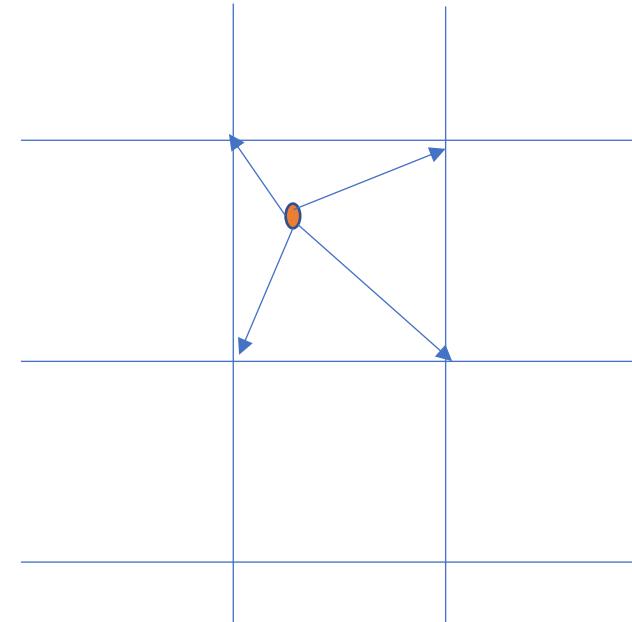
- Number of Particles is defined in the input file
- Two input files.
 - GEM_2D.inp
 - GEM_3D.inp

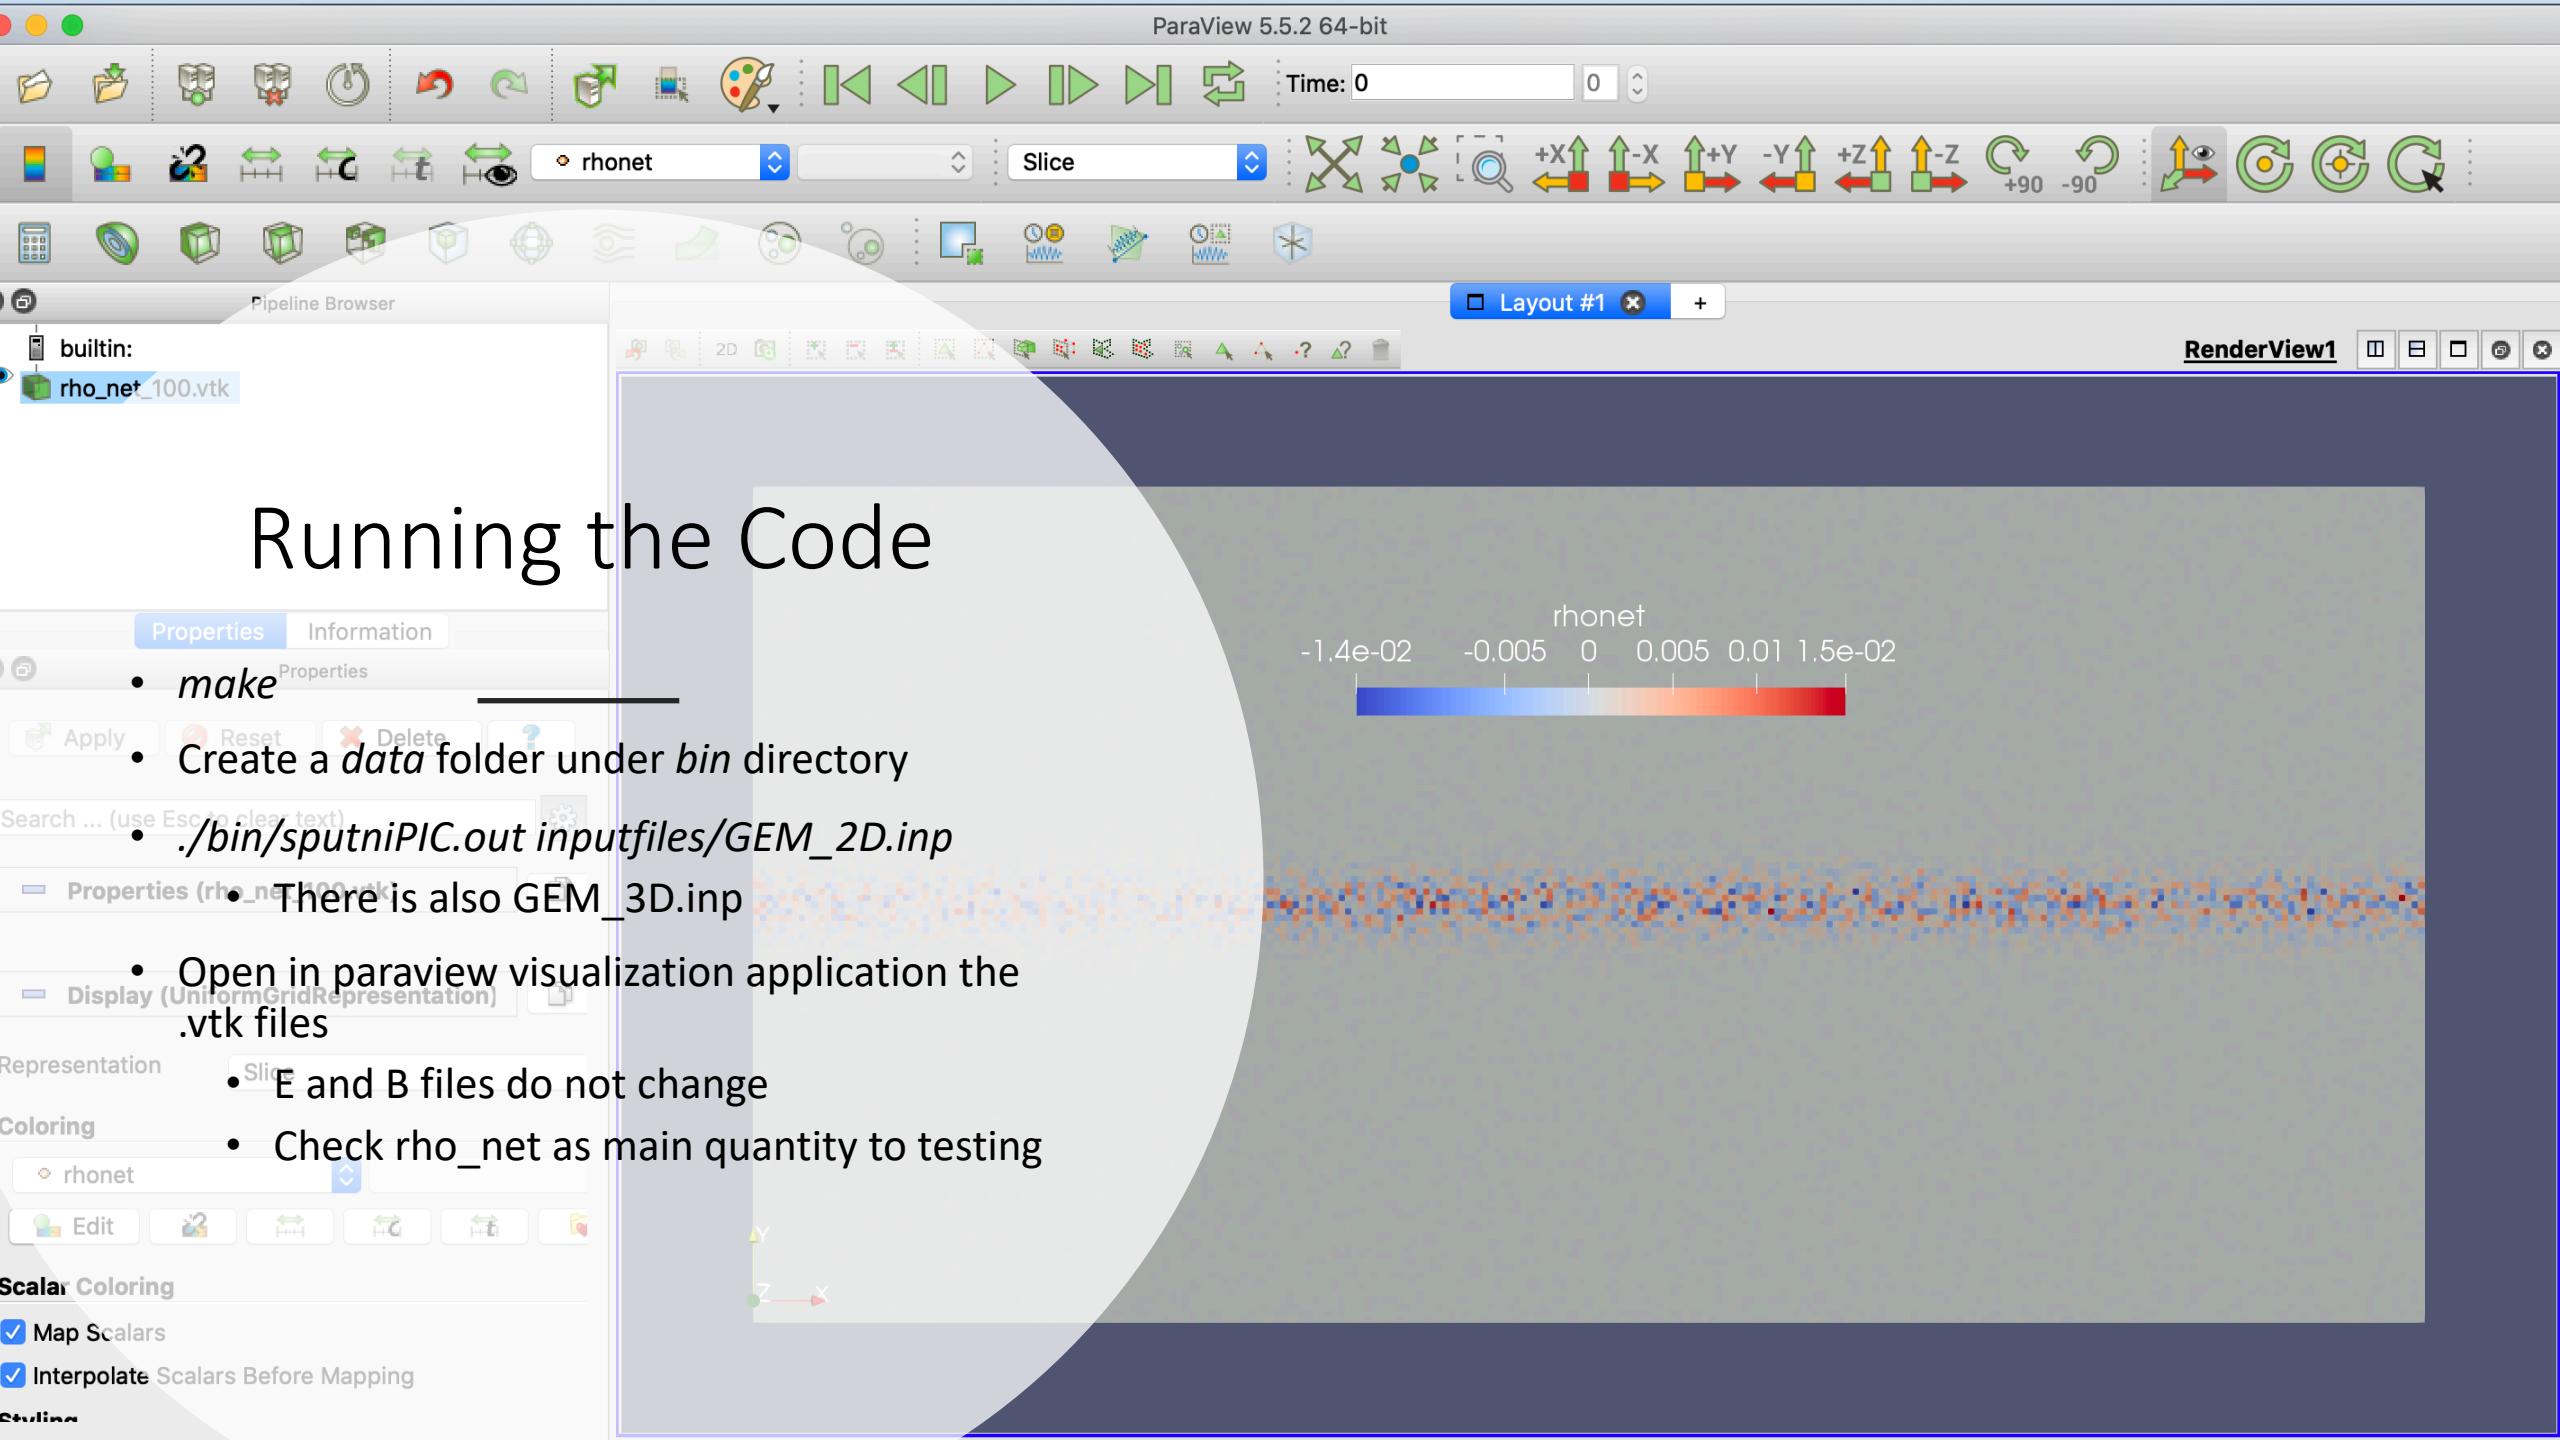
```
13 # %%%%%%%%%%%%%% BOX SIZE %%%%%%%%%%%%%%
14 Lx = 40    # Lx = simulation box length - x direction
15 Ly = 20    # Ly = simulation box length - y direction
16 nxc = 256   # nxc = number of cells - x direction
17 nyc = 128   # nyc = number of cells - y direction
18
19
20 # %%%%%%%%%%%%%% PARTICLES %%%%%%%%%%%%%%
21 #      ns = number of species
22 ns = 4      ← 4 particles species/populations
23
24 # qom = charge to mass ratio for different species */
25 qom = -64.0 1.0 -64 1.0
26                                     ← - = electron, + = proton
27 # Initial density (make sure that plasma is neutral)
28 rhoINIT = 1.0 1.0 0.02 0.02
29
30 # TrackParticleID[species] = 1=true, 0=false --> Assign ID to particles
31 TrackParticleID= 0 0 0 0
32 # npcelx = number of particles per cell - Direction X
33 npcelx = 5 5 5 5
34 # npcely = number of particles per cell - Direction Y */
35 npcely = 5 5 5 5
36 # npcelz = number of particles per cell - Direction Z */
37 npcelz = 5 5 5 5
```

Interpolation Particle to Grid

- At each iteration, 10 quantities defined on the grid nodes (ρ , j_x , j_y , j_z , p_{xx} , p_{xy} , p_{xz} , p_{yy} , p_{yz} and p_{zz}). are deposited from particles on the grid nodes
- In Particles.cu, interpP2G function
- This function need to use atomic operations
 - $\rho[i_x][i_y][i_z] += ...$

atomic





Hints

- The code uses many 3D arrays for E, B, rho, jx, jy, jz, pxx, pxy, pxz, pyy, pyz and pzz
 - In order to move these arrays to GPU you need 1D array
 - Use X_flat quantities
 - Use helper functions in Alloc.h

Detailed Objectives

- Port the particle mover and interpolation functions to GPU with CUDA
- Optimize them
- Measure Performance
- Compare performance with CPU and OpenACC

Check grading criteria



https://kth.instructure.com/courses/12406/pages/grading-criteria?module_item_id=179237

Grade	Code
A	<ul style="list-style-type: none">• Completely offload both the particle mover and interpolation to CUDA, with both steps performed in a single kernel.• Implement a simple particle mover that uses OpenACC.
B	<ul style="list-style-type: none">• Use pinned memory on the host.• Use Stream and async copy
C	<ul style="list-style-type: none">• Implement mini-batches of particles to the particle mover so that the application can process more particles than it fits on the GPU.
D	<ul style="list-style-type: none">• Port part of the interpolation (interpP2G) to CUDA.• All the CUDA memory management is correctly taken care of.
E	<ul style="list-style-type: none">• Port the particle mover (mover_PC) of the code to use GPU, without further optimization.• All the CUDA memory management is correctly taken care of.• All the input files that are used to perform the experiments are provided.

Two Deliverables

- **Design Document.** Due on Dec. 8 2019. P/F
- **Final Report + code** on GitHub. Due on Jan. 14. A-E.
 - This is the final exam
- More instructions on the structure of these documents on the assignment page